

Rethinking ASTE: A Minimalist Tagging Scheme Alongside Contrastive Learning

Anonymous ACL submission

Abstract

Aspect Sentiment Triplet Extraction (ASTE) is a burgeoning subtask of fine-grained sentiment analysis, aiming to extract structured sentiment triplets from unstructured textual data. Existing approaches to ASTE often complicate the task with additional structures or external data. In this research, we propose a novel tagging scheme and employ a contrastive learning approach to mitigate these challenges. The proposed approach demonstrates comparable or superior performance in comparison to state-of-the-art techniques, while featuring a more compact design and reduced computational overhead. Notably, even in the era of Large Language Models (LLMs), our method exhibits superior efficacy compared to GPT 3.5 and GPT 4 in a few-shot learning scenarios. This study also provides valuable insights for the advancement of ASTE techniques within the paradigm of LLMs.

1 Introduction

Aspect Sentiment Triplet Extraction (ASTE) is an emerging fine-grained¹ sentiment analysis task (Pontiki et al., 2014, 2015, 2016) aimed at identifying and extracting structured sentiment triplets (Peng et al., 2020), defined as (Aspect, Opinion, Sentiment), from unstructured text. Specifically, an Aspect term refers to the subject of discussion, an Opinion term provides a qualitative assessment of the Aspect, and Sentiment denotes the overall sentiment polarity, typically taken from a three-level scale (Positive, Neutral, Negative). For instance, consider the sentence: “The battery life is good, but the camera is mediocre.” The ground truth result is {(battery life, good,

¹Generally, sentiment analysis can be based on three levels, namely, document-based, sentence-based, and aspect-based (Jing et al., 2021).




Figure 1: An example for the ASTE task illustrating Aspect terms in purple, Opinion terms with Negative sentiment in blue, and Opinion terms with Positive sentiment in orange.

Positive), (camera, mediocre, Neutral)}. Figure 1 illustrates another example. Recent methods (Wu et al., 2020b; Jing et al., 2021; Chen et al., 2022; Zhang et al., 2022; Liang et al., 2023) commonly utilize Pretrained Language Models (PLMs) to encode input text. The powerful representational capacity of PLMs has greatly advanced the performance in this field, yet there is a tendency to employ complex classification head designs and leverage information enhancement techniques to achieve marginal performance improvements.

In this work, we attribute the current challenges in ASTE to two main factors: 1) *the longstanding overlook of the conical embedding distribution problem* and 2) *imprudent tagging scheme design*. We critically examine the conventional 2D tagging method, commonly known as the table-filling approach, to reassess the efficacy of tagging schemes. Highlighting the critical role of tagging scheme optimization, we delve into what constitutes an ideal scheme for ASTE. We analyze the advantages of the full matrix approach over the half matrix approach and decompose the labels in the full matrix into 1) *location* and 2) *classification*. Thanks to this, we come up with a new tagging scheme with a minimum number of labels to effectively reduce the complexity of training and inference. Moreover, this tagging scheme can be

well aligned with our novel contrastive learning mechanism. To the best of our knowledge, this is the first formal analysis of the tagging scheme to guide a rational design and the first attempt to adopt token-level contrastive learning to improve the PLM representations’ distribution and facilitate the learning process.

The contributions of this work can be summarized as follows:

- We offer the first critical evaluation of the 2D tagging scheme, particularly focusing on the table-filling method. This analysis pioneers in providing a structured framework for the rational design of tagging schemes.
- We introduce a simplified tagging scheme with the least number of label categories to date, integrating a novel token-level contrastive learning approach to enhance PLM representation distribution.
- Our study addresses ASTE challenges in the context of LLMs, developing a tailored in-context learning strategy. Through evaluations on GPT 3.5-Turbo and GPT 4, we establish our method’s superior efficiency and effectiveness.

2 Method

Figure 2 presents our framework. Essentially, our method contributes by two aspects: 1) a *contrastive-learning-enhanced PLM encoder* and 2) a *minimalist tagging scheme*. **Appendix A** Algorithm 1 further delivers a pseudo-code for the training process of our proposed framework.

2.1 The Contrastive-learning-enhanced PLM Encoder

Note that a representation distribution satisfying *alignment* and *uniformity* is linearly separable (Wu et al., 2023) and facilitates the classification. Thereby, contrastive learning boosts represent learning by improving the *alignment* and *uniformity* of the representations (Wu et al., 2023). However, recent investigations indicate that representation distributions in pretrained models often diverge from these expectations. Liang et al. (2021) computed similarities for randomly sampled word pairs, revealing that word embeddings in an Euclidean space cluster within a confined cone, rather than uniformly distributed.

The motivation to adopt contrastive learning is hence to improve the distribution of representations output by PLMs. The core idea lies in that, after fine-tuning on a specific task, the PLM encoder should embed words with similar roles to distribute closer and drive the different ones to be farther.

Given the input sentence \mathcal{S} and pretrained model Encoder , it outputs the hidden word embeddings $\mathcal{H}_{|\mathcal{H}|}$:

$$\mathcal{H}_{|\mathcal{H}|} = \text{Encoder}(\mathcal{S}_{|\mathcal{S}|}), \quad (1)$$

where $\mathcal{H}_{|\mathcal{H}|} = (h_1, h_2, \dots, h_{|\mathcal{H}|})$. Note that, $|\mathcal{H}|$ represents the number of tokens \mathcal{H} , which is not necessarily equal the number of words \mathcal{S} .

Inspired by (Schroff et al., 2015), we 1) take the Euclidean distance as a negative metric on the similarity and 2) introduce a margin d to enforce the gap between two similar hidden word representations, that is, stop pulling \mathcal{H}_i and \mathcal{H}_j closer when there is $\|\mathcal{H}_i - \mathcal{H}_j\|^2 \leq d$ (avoiding similar representations from squeezing too much with each other).

So, the metric of similarity is

$$\text{Sim}_{i,j} = \text{Sim}(\mathcal{H}_i, \mathcal{H}_j) = -\|\mathcal{H}_i - \mathcal{H}_j\|^2 \quad (2)$$

Hence, the similarities $\text{Sim}_{i,j}, i, j \in \{1, 2, \dots, |\mathcal{H}|\}$ forms a similarity matrix $\text{Sim}_{|\mathcal{H}| \times |\mathcal{H}|}$.

To make the representations closer among tokens within the same class and farther between that of different classes, while keeping a margin of d , we can simply maximize $\text{Sim}_{i,j}$ if \mathcal{H}_i and \mathcal{H}_j shares the same class and else minimize $\text{Sim}_{i,j}$.

Defining a matrix $\mathbf{M}_{|\mathcal{H}| \times |\mathcal{H}|}$, it controls whether to pull (closer) or push (farther) between the hidden representation of words, that is, the “Contrastive Mask” used to calculate the contrastive loss. In Figure 2, it is the left-side strict upper-triangle matrix in blue, orange and grey².

²The meaning of colors: The blue cell means that on this cell the row representation is in a different class from the column representation, and the orange cell means that with same class.

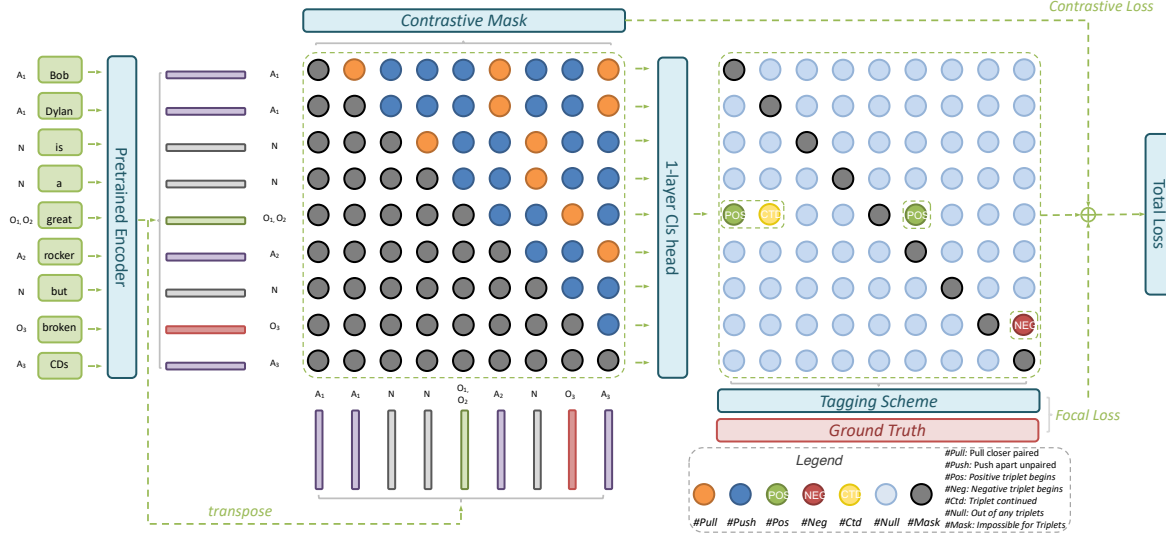


Figure 2: Schematic diagram of our proposed framework which contains: 1) a contrastive learning mechanism that aligns Aspect and Opinion; and 2) a tagging scheme encompassing 5 label classes: NULL, CTD, POS, NEU, and NEG.

$$\begin{aligned}
& \mathcal{L}_{contrastive} \\
&= \sum_{(\mathcal{H}_i, \mathcal{H}_j)^+} \max\{\|\mathcal{H}_i - \mathcal{H}_j\|^2, d\} \\
&- \sum_{(\mathcal{H}_i, \mathcal{H}_j)^-} \min\{\|\mathcal{H}_i - \mathcal{H}_j\|^2, d\} \\
&= \sum_{i=1}^{|\mathcal{H}|} \sum_{j=i+1}^{|\mathcal{H}|} \max\{\text{Sim}(\mathcal{H}_i, \mathcal{H}_j), d\} \cdot \mathbf{M}_{i,j} \\
&= \sum_{i=1}^{|\mathcal{H}|} \sum_{j=1}^{|\mathcal{H}|} (\max\{\text{Sim}_{|\mathcal{H}| \times |\mathcal{H}|}, d\} \circ \mathbf{M}_{|\mathcal{H}| \times |\mathcal{H}|})_{i,j}
\end{aligned} \tag{3}$$

where \circ is a notation for the Hadamard product (Horn, 1990).

Figure 3 depicts the workflow of our contrastive learning strategy and showcases the detail of $\mathbf{M}_{|\mathcal{H}| \times |\mathcal{H}|}$. Note that for the purpose of contrastive learning, the order of each pair of words is redundant, so using the strict upper matrix is enough and we mask the lower triangle.

Noteworthy, a PLM encoder, like BERT, encodes the same word input to different vector representations when the context is different. For example, suppose two sentences, 1) “I really just like the old school plug-in ones”, and 2) “The old school is beautiful”. Explicitly the “school” should be labelled as *Aspect* in the former case whilst as *Opinion* in the latter. For these cases, a well-learned PLM encoder will encode the unique word “school” into different representations according to the context. This characteristic allows

for the direct application of contrastive learning to the hidden word representations without worrying about contradiction.

2.2 The Minimalist Tagging Scheme

Rethinking the 2D tagging scheme:

Lemma 1. Specific to the ASTE task, when we take it as a 2D-labeling problem, we are to 1) find a set of tagging strategies to establish a 1-1 map between each triplet and its corresponding tagging matrix. See the proof in **Appendix Proof 1**.

Lemma 2. In a 2D-tagging for ASTE, at least three basic goals must be met: 1) correctly identifying the (Aspect, Opinion) pairs, 2) correctly classifying the sentiment polarity of the pair based on the context, and 3) avoiding boundary errors, such as *overlapping*³, *confusion*⁴, and *conflict*⁵. See the proof in **Appendix Proof 2**.

Theorem 1. From insight of the above lemmas, it can be concluded that using **enough** (that is, following the 1-1 map properties in Lemma 1, as well as avoiding the issues in Lemma 2) labels will make it a theoretically ensured tagging scheme.

Assumption 1. Ceteris paribus, for a specific classification neural network, the **fewer** the number

³It occurs when one single word belongs to multiple classes in different triplets.

⁴It occurs when there is a lack of location restrictions so that multiple neighbored candidates can not be uniquely distinguished.

⁵It occurs when one single word is composed of multiple tokens, and the predict gives predictions that are not aligned with the word span.




Figure 3: A more detailed illustration of our contrastive learning mechanism.

of target categories, the easier it is for the network to learn. This is an empirical and heuristic assumption, for the reasonable consideration of *Simplification of Decision Boundaries* (Hinton and Salakhutdinov, 2006) and *Enhancement of Training Efficiency* (less parameters).

Combining Theorem 1 and Assumption 1, **fewer yet enough** labels can be heuristically better solution with theoretical guarantee.

With the above knowledge, our tagging scheme employs a full matrix (illustrated as Figure 4) so that rectangular occupations in its cells indicate respective triplets, where each of the rectangles’ row indices correspond to the relative *Aspect* term and the column indices correspond to the *Opinion*. Hereafter, this kind of labels can be taken as a set of “place holder”, which is obviously a 1-1 map meeting Lemma 1.

To further satisfy Lemma 2, we introduce another kind of labels, “sentiment & beginning tag”. This set of labels specializes in recognizing the top-left corner of a “shadowed” area. Meanwhile, it takes a value from the sentiment polarity, i.e. *Positive*, *Neutral*, *Negative*. This tagging is crucial to both *identify the beginning of an triplet* and *label the sentiment polarity*.

Figure 4 shows a comprehensive case of our tagging scheme, in which the left matrix is an appearance of our tagging scheme, and it can be decomposed into two separate components. The middle matrix is the first component, which takes only one tag to locate the up-left beginning of an area, and the second component simply predicts a binary classification to figure out the full area.

Note that, this design benefits the tagging scheme’s decode process. By scanning across the matrix, we only start an examination function when triggered by a beginning label like this, and

then search by row and column until it meets any label except a “continued” (“CTD”), which satisfies Lemma 2.

Thanks to the above design, the ASTE task is well addressed with **ONLY** employing a simple classification head as follows:

$$\begin{aligned} H_1 &= \text{LinearClsHead}(\mathcal{H}_{\text{contrasted}}) \\ H_3 &= \text{LayerNorm}(H_2) \\ \text{Tag}_{\text{pred}} &= \text{GELU}(H_3) \end{aligned} \quad (4)$$

3 Experiments

3.1 Implementation Details

All experiments were conducted on a single RTX 2080 Ti. The best model weight on the development set is saved and then evaluated on the test set. For the PLM encoder, the pretrained weights `bert_base_uncased` and `roberta_base` are downloaded from (Wolf et al., 2020). GPT 3.5-Turbo and GPT 4 are implemented using OpenAI API (OpenAI, 2024). The learning rate is 1×10^{-5} for the PLM encoder, and 1×10^{-3} for the classification head.

3.2 Datasets

We evaluate our method on two canonical ASTE datasets derived from the SemEval Challenges (Pontiki et al., 2014, 2015, 2016). These datasets serve as benchmarks in the majority of Aspect-based Sentiment Analysis (ABSA) research. The first dataset, denoted as \mathcal{D}_1 , is the Aspect-oriented Fine-grained Opinion Extraction (AFOE) dataset introduced by (Wu et al., 2020a). The second dataset, denoted as \mathcal{D}_2 , is a refined version by (Xu et al., 2020), building upon the work of (Peng et al., 2020). Further details are provided in Table 5.




Figure 4: Decomposition of the tagging scheme into two components: 1) a beginning mark matrix with sentiment labels; and 2) a placeholder matrix denoting regions of triplets with “1”s and default regions with “0”s. Remember that each row is taken as candidates for an *Aspect* and each column is taken as candidates for an *Opinion*. Naturally, each cell in the square matrix can be seen as an ordered pair for a unique candidate of $\langle \text{Aspect}, \text{Opinion} \rangle$. When we simply sum the two components up, we have the left-hand tagging scheme in Figure 4, where the “Sentiment & Beginning Tag” is like a trigger (just like you click your mouse), and the “Place Holder” is like a “continued shift” (continue to hold and drag the mouse to the downright).

3.3 Baselines

We evaluate our method against various techniques including pipeline, sequence-labeling, seq2seq, table-filling and LLM-based approaches. Detailed descriptions for each method can be found in the **Appendix Table 8**.

3.4 Performance on ASTE Task

We evaluate ASTE performance using the widely accepted (Precision, Recall, F1) metrics. The result on dataset \mathcal{D}_2 can be found in Table 1 and on \mathcal{D}_1 is presented in **Appendix Table 7**. The best results are indicated in bold, while the second best results are underlined. Our proposed method consistently achieves state-of-the-art performance or ranks second across all evaluated cases.

Significantly, on dataset \mathcal{D}_1 , proposed method achieves a substantial improvement of 3.08% in F1 score on the 14Lap subset. This improvement is particularly noteworthy considering that the best score on this dataset is the lowest among all the datasets, showcasing our ability to effectively handle challenging instances. Moreover, on the 14Res subset, our F1 score surpasses 76.00+, which, to the best of our knowledge, is the highest reported performance.

Turning to dataset \mathcal{D}_2 , our method outperforms all state-of-the-art approaches by more than 1 percentage point on the 14Res, 14Lap, and 16Res subsets. Only on the 16Res subset, the BDTF method (Zhang et al., 2022) achieves a slightly better performance.

On both datasets, our approach overwhelms GPTs’ substantially, indicating that our proposed method keeps advancement in the LLM era. **Appendix Table 9** showcases more interesting facts of GPT’s performance.

We highlight that we evaluate our method against LLM-based approaches in zero-shot and few-shots cases, which could be blamed for not adopting a full parameter fine-tuning. Nevertheless, our evaluation of GPT models on ASTE tasks stems from curiosity about the processing capabilities of large language models, which is believed to be a common interest in the ASTE community, where the full parameter fine-tuning may result in catastrophic forgetting (Lin et al., 2024) (which is unaffordable in a specific task like ASTE) and is left for future research efforts.

3.5 Performance on Other ABSA Tasks

Our method can also effectively handle other ABSA subtasks, including Aspect Extraction (AE), Opinion Extraction (OE), and Aspect Opinion Pair Extraction (AOPE). AE aims to extract all the (*Aspect*) terms, OE aims to extract all the *Opinion* terms, and AOPE aims to extract all the (*Aspect*, *Opinion*) pairs from raw text. The results for these tasks are presented in Table 6 in the **Appendix**, where our method consistently achieves best F1-scores across nearly all tasks.

Methods	14Res			14Lap			15Res			16Res		
	P	R	F1	P	R	F1	P	R	F1	P	R	F1
Pipeline												
Two-stage ^b (Peng et al., 2020)	43.24	63.66	51.46	37.38	50.38	42.87	48.07	57.51	52.32	46.96	64.24	54.21
Li-unified-R+PD ^d (Peng et al., 2020)	40.56	44.28	42.34	41.04	67.35	51.00	44.72	51.39	47.82	37.33	54.51	44.31
Sequence-tagging												
Span-BART (Yan et al., 2021)	65.52	64.99	65.25	61.41	56.19	58.69	59.14	59.38	59.26	66.60	68.68	67.62
JET (Xu et al., 2020)	70.56	55.94	62.40	55.39	47.33	51.04	64.45	51.96	57.53	70.42	58.37	63.83
Seq2seq												
Dual-MRC (Mao et al., 2021)	71.55	69.14	70.32	57.39	53.88	55.58	63.78	51.87	57.21	68.60	66.24	67.40
BMRC [†] (Chen et al., 2021a)	72.17	65.43	68.64	65.91	52.15	58.18	62.48	55.55	58.79	69.87	65.68	67.35
COM-MRC (Zhai et al., 2022)	75.46	68.91	72.01	62.35	58.16	60.17	68.35	61.24	64.53	71.55	71.59	71.57
Triple-MRC (Zou et al., 2024)	-	-	72.45	-	-	60.72	-	-	62.86	-	-	68.65
Table-filling												
GTS (Wu et al., 2020a)	67.76	67.29	67.50	57.82	51.32	54.36	62.59	57.94	60.15	66.08	66.91	67.93
Double-encoder (Jing et al., 2021)	67.95	71.23	69.55	62.12	<u>56.38</u>	59.11	58.55	60.00	59.27	70.65	70.23	70.44
EMC-GCN (Chen et al., 2022)	71.21	72.39	71.78	61.70	56.26	58.81	61.54	62.47	61.93	65.62	71.30	68.33
BDTF (Zhang et al., 2022)	75.53	<u>73.24</u>	<u>74.35</u>	68.94	55.97	61.74	68.76	<u>63.71</u>	66.12	71.44	<u>73.13</u>	72.27
STAGE-1D (Liang et al., 2023)	79.54	68.47	73.58	<u>71.48</u>	53.97	61.49	72.05	58.23	64.37	78.38	69.10	<u>73.45</u>
STAGE-2D (Liang et al., 2023)	78.51	69.3	73.61	70.56	55.16	61.88	<u>72.33</u>	58.93	64.94	<u>77.67</u>	68.44	72.75
STAGE-3D (Liang et al., 2023)	<u>78.58</u>	69.58	73.76	71.98	53.86	61.58	73.63	57.9	64.79	76.67	70.12	73.24
DGCNAP (Li et al., 2023)	72.90	68.69	70.72	62.02	53.79	57.57	62.23	60.21	61.19	69.75	69.44	69.58
LLM-based												
GPT 3.5 zero-shot	44.88	55.13	49.48	30.04	41.04	34.69	36.02	53.40	43.02	39.92	57.78	47.22
GPT 3.5 few-shots	52.36	54.63	53.47	29.91	36.04	32.69	45.48	61.44	52.01	49.50	67.12	56.98
GPT 4 zero-shot	32.99	38.13	35.37	17.81	22.55	19.90	27.85	37.73	32.05	32.17	43.00	36.80
GPT 4 few-shots	47.25	49.20	48.20	26.04	33.64	29.35	39.94	51.13	44.85	43.72	54.86	48.66
Ours												
ContrASTE	76.1	75.08	75.59	66.82	60.68	63.61	66.50	63.86	<u>65.15</u>	75.52	74.14	74.83

Table 1: Experimental results on \mathcal{D}_2 (Xu et al., 2020). The best results are highlighted in bold, while the second best results are underlined.

Models	\mathcal{D}_1				\mathcal{D}_2			
	14Res	14Lap	15Res	16Res	14Res	14Lap	15Res	16Res
ContrASTE	76.00	64.07	65.43	71.80	75.59	63.61	65.15	74.83
w/o. RoBERTa	74.12	63.18	62.95	69.41	72.66	62.15	63.25	70.71
ΔF_1	-1.88	-0.89	-2.48	-2.39	-2.93	-1.46	-1.90	-4.12
w/o. contr	72.61	61.94	58.14	68.16	71.72	61.49	58.11	68.03
ΔF_1	-3.39	-2.13	-7.29	-3.64	-3.87	-2.12	-7.04	-6.80
w/o. tag	67.78	54.98	60.75	62.62	65.83	54.98	58.73	67.63
ΔF_1	-8.22	-9.09	-4.68	-9.18	-9.76	-8.63	-6.42	-7.20

Table 2: Ablation study on F1, where “w/o. RoBERTa” denotes “Replace RoBERTa with bert-base-uncased”, “w/o. contr” denotes without the contrastive learning mechanism, and “w/o. tag” denotes “replace our tagging scheme with a baseline”.

4 Analysis

4.1 Ablation Study

Encoder. In the ablation experiment part, by replacing RoBERTa by BERT, The results obtained declined slightly while still yield most other methods.

Contrastive Learning. First, we deactivating the contrastive mechanism in our method (denoting “w/o. contr”) by setting the coefficient of the contrastive loss to 0. The results in Table 2 illustrate a significant F1-score decrease of $0.22 \sim 2.93\%$.

Tagging Scheme. Third, we substitute our pro-

posed scheme with the conventional GTS tagging scheme (Wu et al., 2020a), which results in a significant performance decline (Table 2) by $1.46 \sim 6.14\%$. This indicates that the contrastive learning methods, within our framework, is of strong reliance on an appropriate tagging scheme. This reinforces the effectiveness of our straightforward yet impactful tagging scheme.

4.2 Effect of Contrastive Learning

In Appendix Figure 5 an example is shown of how contrastive learning has improved the representation, where the left subplot is the tokens output by RoBERTa without contrastive learning, and the right one is with contrastive learning. Note that the Principal Component Analysis (PCA) (Maćkiewicz and Ratajczak, 1993) is adopted to reduce the dimensions of the vectors to 2.

4.3 Efficiency Analysis

Table 3 shows an efficiency analysis, where our method is in significantly higher efficiency. While our approach demonstrates efficient management of memory and parameter utilization, surpassing other methods in terms of runtime performance, it

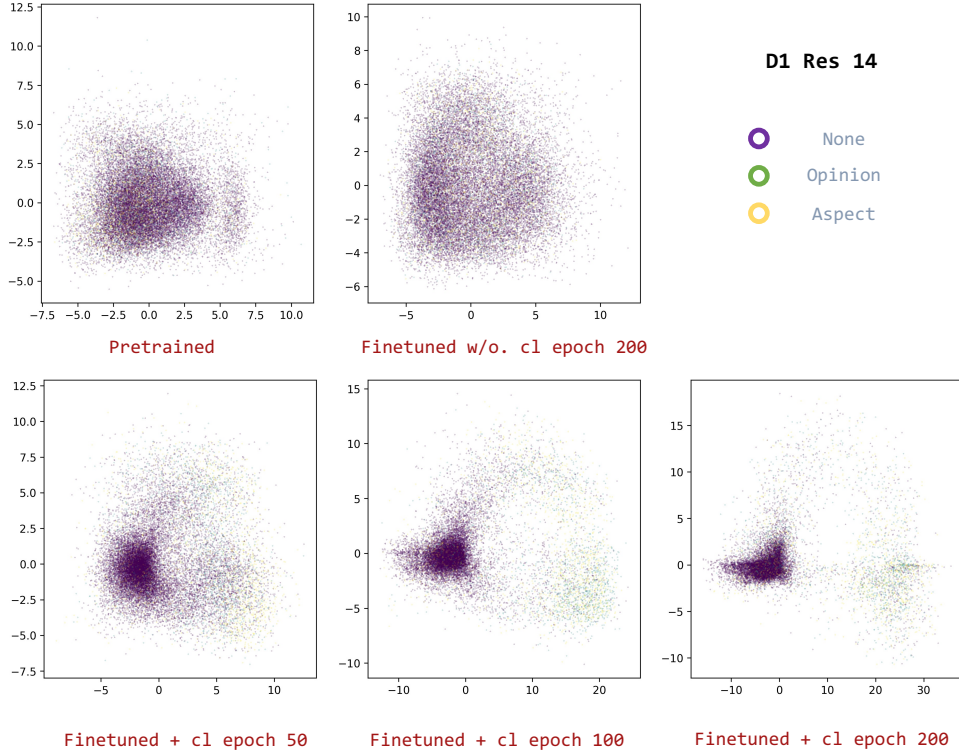


Figure 5: A plot of the hidden word representation, where the dimension is reduced to 2 for convenience of display. “Pretrained” means the model with official released version of weights. “Finetuned” means the model is a finetuned version on ASTE task for certain epochs. “w/o. cl” means the model is trained without contrastive learning loss. “+ cl” means the model is trained with contrastive learning. All the plotted results are from experiment carried on \mathcal{D}_1 14Res.

Model	Memory	Num Params	Epoch Time	Inf Time	F1(%)	Device
Span-ASTE	3.173GB	-	108s	-	71.62	Tesla v100
BDTF	8.103GB	-	135s	-	74.73	Tesla v100
GPT 3.5-Turbo	> 80GB [‡]	175B [†]	-	0.83s	49.48	OpenAI API
GPT 4	> 80GB [‡]	1760B [†]	-	1.56s	35.37	OpenAI API
Ours	7.11GB	0.12B	10s	0.01s	76.00	2080 Ti

Table 3: Efficiency Analysis, where [†] is evaluated by (Gao, 2021) and later confirmed by OpenAI (Wikipedia, 2024), [‡] is estimated by (Schreiner, 2023), and [‡] is reported by (Wikipedia, 2024).

Method	Num Tags	Linguistic Features	Half/Full Matrix
GTS	6	None	Half
Double-encoder	9	None	Half
EMC-GCN	10	4 Groups	Full
BDTF	$2 \times 2 \times 3$	None	Half
STAGE	$2 \times 2 \times 4$	None	Half
DGCNAP	6	POS-tagging	Half
Ours	5	None	Full

Table 4: Tagging Scheme Comparisons

is crucial to note the unsatisfactory performance of LLM-based methods, as demonstrated by Table 7, Table 1 and Table 9. This observation highlights that employing a general LLM with a large number of parameters does not yield desirable results in specific tasks like ASTE, even with the assistance of few-shot memorizing processes. Moreover, the incorporation of LLM introduces a significant computational resource overhead in general application scenarios. Although fine-tuning the parameters of LLM itself may offer some improvement, there exists a risk of collapse forgetting. Furthermore, Table 4 provides a comparative analysis of tagging schemes. Our method achieves its objective without the need for additional linguistic information and utilizes fewer tags.

5 Related Work

5.1 ASTE Paradigms

Peng et al. (2020) proposes a pipeline method to bifurcate ASTE tasks into two stages, extracting (Aspect, opinion) pairs initially and predicting sentiment polarity subsequently. However,

the error propagation hampers pipeline methods, rendering them vulnerable to end-to-end counterparts. End-to-end strategies, under the sequence-labeling approach, treat ASTE as a 1D “B-I-O” tagging scheme. ET (Xu et al., 2020) introduces a position-aware tagging scheme with a conditional random field (CRF) module, effectively addressing span overlapping issues. Recent advances in the end-to-end paradigm delicately grasp the peculiarity of ASTE tasks and come up with a proficient 2D table-filling tagging scheme. Other researches treat ASTE as generation problem, and develops seq2seq methods such as machine reading comprehension (Zhai et al., 2022; Mao et al., 2021; Zou et al., 2024; Chen et al., 2021b).

5.2 Tagging Schemes

Wu et al. (2020a) pioneer the adoption of a grid tagging scheme (GTS) for ASTE, yielding substantial performance gains. While subsequent research refines and enhances GTS, it is not devoid of drawbacks. Instances involving multi-word *Aspect/Opinion* constructs risk *relation inconsistency* and *boundary insensitivity* (Zhang et al., 2022). To overcome these, BDTF (Zhang et al., 2022) designs a boundary-driven tagging scheme, effectively reducing boundary prediction errors. Alternative research augments GTS by integrating external semantic information as structured knowledge into their models. S³E² (Chen et al., 2021b) retains the GTS tagging scheme while introducing novel semantic and syntactic enhancement modules between word embedding outputs and the tagging scheme. EMGCN (Chen et al., 2022) offers a distinct perspective, incorporating external knowledge from four aspects, namely, Part-of-Speech Combination, Syntactic Dependency Type, Tree-based Distance, and Relative Position Distance through an exogenous hard-encoding strategy. SyMux (Fei et al., 2022) contributes a unified tagging scheme capable of all ABSA subtasks synthesizing insights from incorporating GCN, syntax encoder, and representation multiplexing.

5.3 Contrastive Learning

While contrastive learning has gained popularity in diverse NLP domains (Wu et al., 2020b; Giorgi et al., 2021; Gao et al., 2021; Zhang et al., 2021), its application to ASTE remains relatively unexplored. Ye et al. (2021) adopts contrastive learning into triplet extraction in a generative fash-

ion. Wang et al. (2022) takes contrastive learning as a data augmentation approach. Yang et al. (2023) proposed an enhancement approach in pairing with two separate encoders.

6 Conclusion

In this work, we have introduced an elegant and efficient framework for ASTE, achieving SOTA performance. Our approach is built upon two effective components: a new tagging scheme and a novel token-level contrastive learning implementation. The ablation study demonstrates the synergy between these components, reducing the need for complex model designs and external information enhancements.

7 Limitations & Future Work

Our work focused on ASTE problem for English, and can not ensure that our framework will work so well in other languages. However, our conventional 2D tagging method doesn’t care much about grammar or other rules of a language. We believe that our framework will apply in other languages.

Our study is purely based on the table-filling paradigm of ASTE approaches. In future work, it is worth exploring a combination between our method with other paradigms, such as seq2seq approaches.

Finally, our framework may not be the best. It leaves potential to further investigate various classification head strategies.

8 Ethics Statement

In all our experiments, we employed datasets that are widely accepted and extensively referenced within the academic community. These datasets primarily focus on reviews of products and services on e-commerce platforms, which inherently possess a lower risk of containing offensive content. We have made every effort to scrutinize the data for potential biases against gender, race, and marginalized groups.

Despite our precautions, it is important to note that our model might still generate sentiment assessments that could be perceived as offensive, particularly if deployed in inappropriate contexts, such as evaluating statements related to ethical or moral issues. In such cases, we reserve the right to limit or modify the use of our technology to prevent misuse.

References

- Hao Chen, Zepeng Zhai, Fangxiang Feng, Ruifan Li, and Xiaojie Wang. 2022. Enhanced multi-channel graph convolutional network for aspect sentiment triplet extraction. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2974–2985.
- Shaowei Chen, Yu Wang, Jie Liu, and Yuelin Wang. 2021a. Bidirectional machine reading comprehension for aspect sentiment triplet extraction. In *Proceedings of the AAAI conference on artificial intelligence*, volume 35, pages 12666–12674.
- Zhexue Chen, Hong Huang, Bang Liu, Xuanhua Shi, and Hai Jin. 2021b. Semantic and syntactic enhanced aspect sentiment triplet extraction. *arXiv preprint arXiv:2106.03315*.
- Hongliang Dai and Yangqiu Song. 2019. Neural aspect and opinion term extraction with mined rules as weak supervision. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 5268–5277.
- Hao Fei, Fei Li, Chenliang Li, Shengqiong Wu, Jingye Li, and Donghong Ji. 2022. Inheriting the wisdom of predecessors: A multiplex cascade framework for unified aspect-based sentiment analysis. In *Proceedings of the Thirty-First International Joint Conference on Artificial Intelligence, IJCAI*, pages 4096–4103.
- Leo Gao. 2021. On the sizes of openai api models. <https://blog.eleuther.ai/gpt3-model-sizes/>. Retrieved November 23, 2023.
- Tianyu Gao, Xingcheng Yao, and Danqi Chen. 2021. Simcse: Simple contrastive learning of sentence embeddings. In *2021 Conference on Empirical Methods in Natural Language Processing, EMNLP 2021*, pages 6894–6910. Association for Computational Linguistics (ACL).
- John Giorgi, Osvald Nitski, Bo Wang, and Gary Bader. 2021. Declutr: Deep contrastive learning for unsupervised textual representations. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 879–895.
- Geoffrey E. Hinton and Ruslan R. Salakhutdinov. 2006. Reducing the dimensionality of data with neural networks. *Science*, 313(5786):504–507.
- Roger A Horn. 1990. The hadamard product. In *Proc. Symp. Appl. Math.*, volume 40, pages 87–169.
- Hongjiang Jing, Zuchao Li, Hai Zhao, and Shu Jiang. 2021. Seeking common but distinguishing difference, a joint aspect-based sentiment analysis model. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 3910–3922.
- Yanbo Li, Qing He, and Damin Zhang. 2023. Dual graph convolutional networks integrating affective knowledge and position information for aspect sentiment triplet extraction. *Frontiers in Neurorobotics*, 17.
- Shuo Liang, Wei Wei, Xian-Ling Mao, Yuanyuan Fu, Rui Fang, and Danyang Chen. 2023. Stage: span tagging and greedy inference scheme for aspect sentiment triplet extraction. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37, pages 13174–13182.
- Yuxin Liang, Rui Cao, Jie Zheng, Jie Ren, and Ling Gao. 2021. Learning to remove: Towards isotropic pre-trained bert embedding. In *Artificial Neural Networks and Machine Learning–ICANN 2021: 30th International Conference on Artificial Neural Networks, Bratislava, Slovakia, September 14–17, 2021, Proceedings, Part V 30*, pages 448–459. Springer.
- Yong Lin, Hangyu Lin, Wei Xiong, Shizhe Diao, Jianmeng Liu, Jipeng Zhang, Rui Pan, Haoxiang Wang, Wenbin Hu, Hanning Zhang, Hanze Dong, Renjie Pi, Han Zhao, Nan Jiang, Heng Ji, Yuan Yao, and Tong Zhang. 2024. [Mitigating the alignment tax of rlhf](#).
- Andrzej Maćkiewicz and Waldemar Ratajczak. 1993. Principal components analysis (pca). *Computers & Geosciences*, 19(3):303–342.
- Yue Mao, Yi Shen, Chao Yu, and Longjun Cai. 2021. A joint training dual-mrc framework for aspect based sentiment analysis. In *Proceedings of the AAAI conference on artificial intelligence*, volume 35, pages 13543–13551.
- OpenAI. 2024. OpenAI API. <https://platform.openai.com/docs>. Accessed: 2024-04-14.
- Haiyun Peng, Lu Xu, Lidong Bing, Fei Huang, Wei Lu, and Luo Si. 2020. Knowing what, how and why: A near complete solution for aspect-based sentiment analysis. In *Proceedings of the AAAI conference on artificial intelligence*, volume 34, pages 8600–8607.
- Eleni Pontiki, Dimitra Hadjipavlou-Litina, Konstantinos Litinas, and George Geromichalos. 2014. Novel cinnamic acid derivatives as antioxidant and anti-cancer agents: Design, synthesis and modeling studies. *Molecules*, 19(7):9655–9674.
- Maria Pontiki, Dimitrios Galanis, Harris Papageorgiou, Suresh Manandhar, and Ion Androutsopoulos. 2015. Semeval-2015 task 12: Aspect based sentiment analysis. In *Proceedings of the 9th international workshop on semantic evaluation (SemEval 2015)*, pages 486–495.

- Maria Pontiki, Dimitris Galanis, Haris Papageorgiou, Ion Androutsopoulos, Suresh Manandhar, Mohammed AL-Smadi, Mahmoud Al-Ayyoub, Yanyan Zhao, Bing Qin, Orphée De Clercq, et al. 2016. Semeval-2016 task 5: Aspect based sentiment analysis. In *ProWorkshop on Semantic Evaluation (SemEval-2016)*, pages 19–30. Association for Computational Linguistics.
- Maximilian Schreiner. 2023. Gpt-4 architecture, datasets, costs and more leaked. THE DECODER. Archived from the original on July 12, 2023. Retrieved July 12, 2023.
- Florian Schroff, Dmitry Kalenichenko, and James Philbin. 2015. Facenet: A unified embedding for face recognition and clustering. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 815–823.
- Bing Wang, Liang Ding, Qihuang Zhong, Ximing Li, and Dacheng Tao. 2022. A contrastive cross-channel data augmentation framework for aspect-based sentiment analysis. In *Proceedings of the 29th International Conference on Computational Linguistics*, pages 6691–6704.
- Wenya Wang, Sinno Jialin Pan, Daniel Dahlmeier, and Xiaokui Xiao. 2017. Coupled multi-layer attentions for co-extraction of aspect and opinion terms. In *Proceedings of the Thirty-First AAAI Conference on Artificial Intelligence*, pages 3316–3322.
- Wikipedia. 2024. Gpt-3 - wikipedia. <https://en.wikipedia.org/wiki/GPT-3>. (Accessed on 04/14/2024).
- Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, et al. 2020. Transformers: State-of-the-art natural language processing. In *Proceedings of the 2020 conference on empirical methods in natural language processing: system demonstrations*, pages 38–45.
- Jing Wu, Jennifer Hobbs, and Naira Hovakimyan. 2023. Hallucination improves the performance of unsupervised visual representation learning. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 16132–16143.
- Zhen Wu, Chengcan Ying, Fei Zhao, Zhifang Fan, Xinyu Dai, and Rui Xia. 2020a. Grid tagging scheme for aspect-oriented fine-grained opinion extraction. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 2576–2585.
- Zhuofeng Wu, Sinong Wang, Jiatao Gu, Madian Khabsa, Fei Sun, and Hao Ma. 2020b. Clear: Contrastive learning for sentence representation. *arXiv preprint arXiv:2012.15466*.
- Lu Xu, Hao Li, Wei Lu, and Lidong Bing. 2020. Position-aware tagging for aspect sentiment triplet extraction. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 2339–2349.
- Hang Yan, Junqi Dai, Tuo Ji, Xipeng Qiu, and Zheng Zhang. 2021. A unified generative framework for aspect-based sentiment analysis. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 2416–2429.
- Fan Yang, Mian Zhang, Gongzhen Hu, and Xiabing Zhou. 2023. A pairing enhancement approach for aspect sentiment triplet extraction. *arXiv preprint arXiv:2306.10042*.
- Hongbin Ye, Ningyu Zhang, Shumin Deng, Mosh Chen, Chuanqi Tan, Fei Huang, and Huajun Chen. 2021. Contrastive triple extraction with generative transformer. In *Proceedings of the AAAI conference on artificial intelligence*, volume 35, pages 14257–14265.
- Zepeng Zhai, Hao Chen, Fangxiang Feng, Ruifan Li, and Xiaojie Wang. 2022. Com-mrc: A context-masked machine reading comprehension framework for aspect sentiment triplet extraction. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 3230–3241.
- Chen Zhang, Qiuchi Li, Dawei Song, and Benyou Wang. 2020. A multi-task learning framework for opinion triplet extraction. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 819–828.
- Dejiao Zhang, Feng Nan, Xiaokai Wei, Shang-Wen Li, Henghui Zhu, Kathleen Mckeown, Ramesh Nallapati, Andrew O Arnold, and Bing Xiang. 2021. Supporting clustering with contrastive learning. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 5419–5430.
- Yice Zhang, Yifan Yang, Yihui Li, Bin Liang, Shiwei Chen, Yixue Dang, Min Yang, and Ruifeng Xu. 2022. Boundary-driven table-filling for aspect sentiment triplet extraction. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 6485–6498.
- Wang Zou, Wubo Zhang, Wenhuan Wu, and Zhuoyan Tian. 2024. A multi-task shared cascade learning for aspect sentiment triplet extraction using bert-mrc. *Cognitive Computation*, pages 1–18.

A Appendix

Proof 1:

Let:

- S be a sentence with n tokens.
- M be an $n \times n$ tagging matrix for S , where each entry $M[i][j]$ can hold a label.
- $T_k = (A_k, O_k, S_k)$ be a sentiment triplet consisting of an aspect term A_k , an opinion term O_k , and a sentiment S_k .

Tagging Strategy If A_k starts at position i and O_k starts at position j , then $M[i][j]$ is tagged with a unique label L_k that encodes S_k . This label L_k uniquely identifies the triplet T_k , ensuring that no other entry $M[i'][j']$ with $(i', j') \neq (i, j)$ carries the same label unless it refers to the same sentiment context.

Define $L_k = \text{"start of triplet"}T_k$ with sentiment S_k

Proof of One-to-One Mapping

- **Injectivity:** Each L_k uniquely identifies a triplet T_k . If $M[i][j] = M[i'][j'] = L_k$, then by definition, $(i, j) = (i', j')$ and T_k is the same.
- **Surjectivity:** Each triplet T_k can be uniquely located and identified by its label L_k in matrix M , where no two distinct triplets have the same label at the same matrix position.

Conclusion The tagging scheme ensures that each sentiment triplet T_k is uniquely mapped to a specific label in the matrix M , and each label in M uniquely refers back to a specific triplet T_k . This guarantees a one-to-one correspondence between the triplets and their tagging matrix representations, fulfilling the conditions required by Lemma 1 for an effective and efficient ASTE process.

Proof 2:

For the ASTE task, considered as a 2D-labeling problem, it is necessary to ensure three fundamental goals are met:

Definitions

- S be a sentence with n tokens.
- M be an $n \times n$ tagging matrix for S , where each entry $M[i][j]$ can hold a label indicating a component of a sentiment triplet.

- $T_k = (A_k, O_k, S_k)$ be a sentiment triplet consisting of an aspect term A_k , an opinion term O_k , and a sentiment S_k .

Goals

1. **Correct Identification of Pairs:** Ensure that each (Aspect, Opinion) pair is correctly identified in the tagging matrix M .
2. **Classification of Sentiment Polarity:** Accurately classify the sentiment polarity S_k for each (Aspect, Opinion) pair.
3. **Avoidance of Boundary Errors:** Prevent boundary errors such as overlapping and confusion in the tagging matrix M .

Proof Using Contraposition

1. **Assuming Incorrect Identification:** Assume that some (Aspect, Opinion) pairs are incorrectly identified in M . This would mean that there exists at least one pair (i, j) where $M[i][j]$ does not represent the actual (Aspect, Opinion) relationship in S . This misrepresentation leads to incorrect sentiment analysis results, which contradicts the requirement of the task to provide accurate sentiment analysis, thereby proving that our identification must be correct.
2. **Assuming Incorrect Classification:** Assume the sentiment polarity S_k is incorrectly classified in M . This would imply that the sentiment associated with an (Aspect, Opinion) pair is wrong, leading to a sentiment analysis that does not reflect the true sentiment of the text. Given that the primary goal of ASTE is to accurately identify sentiments, this assumption leads to a contradiction, thereby establishing that our classification must be accurate.
3. **Assuming Existence of Boundary Errors:** Assume boundary errors such as overlaps or confusion occur in M . Such errors would prevent the clear identification and classification of sentiment triplets, leading to incorrect or ambiguous extraction outcomes. This would undermine the integrity and usability of the ASTE process, contradicting the task's need for precise extraction mechanisms. Hence, we prove that boundary errors must be effectively managed.

Conclusion The contraposition approach solidifies that the tagging strategy for ASTE in

a 2D labeling framework successfully achieves the correct identification of pairs, accurate classification of sentiment, and effective management of boundary errors, as any failure in these aspects leads to contradictions with the task requirements.

Algorithm 1 Workflow of our framework.

Modules:

Input:

Raw sentences: $\mathcal{S}_{|\mathcal{S}|}$;

Ground truth triplets: $\mathcal{T}_{|\mathcal{T}|}^{gt}$, where

$\mathcal{T}_k = (A_k, O_k, S_k)$, $k \in \{1, 2, \dots, |\mathcal{T}|\}$;

classes of contrasted labels: \mathcal{C} .

Output:

Predicted Triplets: $\mathcal{T}_{|\mathcal{T}|}^{pred}$;

Metric: *Precision, Recall, F1*.

Algorithm:

Repeat for N epochs:

1: Hidden word representation:

$\mathcal{H}_{|\mathcal{H}|} = \text{PLMsEncoder}(\mathcal{S}_{|\mathcal{S}|})$;

2: Tensor Operations:

$\mathcal{H}_{|\mathcal{H}| \times |\mathcal{H}|} = \text{expand}(\mathcal{H}_{|\mathcal{H}|})$,

$\mathcal{H}_{|\mathcal{H}| \times |\mathcal{H}|}^T = \mathcal{H}_{|\mathcal{H}| \times |\mathcal{H}|} \cdot \text{transpose}()$;

3: Similarity matrix:

$\mathbf{Sim}_{|\mathcal{H}| \times |\mathcal{H}|} = -(\mathcal{H}_{|\mathcal{H}| \times |\mathcal{H}|} - \mathcal{H}_{|\mathcal{H}| \times |\mathcal{H}|}^T) \circ (\mathcal{H}_{|\mathcal{H}| \times |\mathcal{H}|} - \mathcal{H}_{|\mathcal{H}| \times |\mathcal{H}|}^T)$, where $\mathbf{Sim}_{i,j} = -\|\mathcal{H}_i - \mathcal{H}_j\|^2$, and \circ denotes the Hadamard product;

4: Contrastive Mask matrix: $\mathbf{M}_{|\mathcal{H}| \times |\mathcal{H}|}$, where $\mathbf{M}_{i,j} = 1$ if $\mathcal{H}_i, \text{mathcal{H}}_j \in \mathcal{C}_p$, $p \in 1, 2, 3$ else -1 ;

5: Contrastive loss:

$\mathcal{L}_{\text{contrastive}} = \sum_{i=1}^{|\mathcal{H}|} \sum_{j=1}^{|\mathcal{H}|} (\mathbf{Sim}_{|\mathcal{H}| \times |\mathcal{H}|} \circ \mathbf{M}_{|\mathcal{H}| \times |\mathcal{H}|})_{i,j}$;

6: Predicted tagging matrix:

$\mathbf{Tag}_{|\mathcal{H}| \times |\mathcal{H}|}^{pred} = \text{ClsHead}(\mathcal{H}_{|\mathcal{H}| \times |\mathcal{H}|}, \mathcal{H}_{|\mathcal{H}| \times |\mathcal{H}|}^T)$;

7: Focal loss:

$\mathcal{L}_{\text{focal}} = \text{FocalLoss}(\mathbf{Tag}_{|\mathcal{H}| \times |\mathcal{H}|}^{pred}, \mathbf{Tag}_{|\mathcal{H}| \times |\mathcal{H}|}^{gt})$;

8: Weighted Loss: $\mathcal{L} = \mathcal{L}_{\text{focal}} + \alpha \mathcal{L}_{\text{contrastive}}$.

9: Backward propagation.

Predicted triplets:

$\mathcal{T}_{|\mathcal{T}|}^{pred} = \text{TaggingDecoder}(\mathbf{Tag}_{|\mathcal{H}| \times |\mathcal{H}|}^{pred})$

Metric:

$\text{Precision, Recall, F1} = \text{Metric}(\mathcal{T}_{|\mathcal{T}|}^{pred}, \mathcal{T}_{|\mathcal{T}|}^{gt})$

Datasets		#S	#A	#O	#S1	#S2	#S3	#T	
14Res	\mathcal{D}_1	Train	1259	1008	849	1456	164	446	2066
		Dev	315	358	321	352	44	93	489
		Test	493	591	433	651	59	141	851
	\mathcal{D}_2	Train	1266	986	844	1692	166	480	2338
		Dev	310	396	307	404	54	119	577
		Test	492	579	437	773	66	155	994
14Lap	\mathcal{D}_1	Train	899	731	693	691	107	466	1264
		Dev	225	303	237	173	42	118	333
		Test	332	411	330	305	62	101	468
	\mathcal{D}_2	Train	906	733	695	817	126	517	1460
		Dev	219	268	237	169	36	141	346
		Test	328	400	329	364	63	116	543
15Res	\mathcal{D}_1	Train	603	585	485	668	24	179	871
		Dev	151	182	161	156	8	41	205
		Test	325	353	307	293	19	124	436
	\mathcal{D}_2	Train	605	582	462	783	25	205	1013
		Dev	148	191	183	185	11	53	249
		Test	322	347	310	317	25	143	485
16Res	\mathcal{D}_1	Train	863	775	602	890	43	280	1213
		Dev	216	270	237	224	8	66	298
		Test	328	342	282	360	25	72	457
	\mathcal{D}_2	Train	857	759	623	1015	50	329	1394
		Dev	210	251	221	252	11	76	339
		Test	326	338	282	407	29	78	514

Table 5: Statistic information of our two experiment datasets: “#S”, “#T”, “#A”, and “#O” denote the numbers of “Sentences”, “Triplets”, “Aspects”, and “Opinions”; “#S1”, “#S2”, #S3” denote the numbers of sentiments “Positive”, “Neutral” and “Negative”, respectively.

Methods	14Res			14Lap			15Res			16Res		
	AE	OE	AOPE	AE	OE	AOPE	AE	OE	AOPE	AE	OE	AOPE
CMLA	81.22	83.07	48.95	78.68	77.95	44.10	76.03	74.67	44.60	74.20	72.20	50.00
RINANTE	81.34	83.33	46.29	77.13	75.34	29.70	73.38	75.40	35.40	72.82	70.45	30.70
Li-unified	81.62	85.26	55.34	78.54	77.55	52.56	74.65	74.25	56.85	73.36	73.87	53.75
GTS	83.82	85.04	75.53	79.52	78.61	65.67	78.22	79.31	67.53	75.80	76.38	74.62
Dual-MRC	86.60	86.22	77.68	80.44	79.90	63.37	75.08	77.52	64.97	76.87	77.90	75.71
ContrASTE (Ours)	86.55	87.04	79.60	82.62	83.41	73.23	86.53	83.05	73.87	85.48	87.06	76.29
$\Delta F1$	-0.05	0.82	1.92	2.18	3.51	7.56	8.31	3.74	6.34	8.61	9.16	0.58

Table 6: F1-score performance on other ABSA tasks: AE, OE, and AOPE. The test is implemented on \mathcal{D}_1 . Results of other models are retrieved from (Fei et al., 2022).

Methods	14Res			14Lap			15Res			16Res		
	P	R	F1	P	R	F1	P	R	F1	P	R	F1
Pipeline												
OTE-MTL (Zhang et al., 2020)	-	-	45.05	-	-	59.67	-	-	48.97	-	-	55.83
Li-unified-R+PD [‡] (Peng et al., 2020)	41.44	68.79	51.68	42.25	42.78	42.47	43.34	50.73	46.69	38.19	53.47	44.51
RI-NANTE+ (Dai and Song, 2019)	31.42	39.38	34.95	21.71	18.66	20.07	29.88	30.06	29.97	25.68	22.30	23.87
CMLA+C-GCN [‡] (Wang et al., 2017)	72.22	56.35	63.17	60.69	47.25	53.03	64.31	49.41	55.76	66.61	59.23	62.70
Two-satge [‡] (Peng et al., 2020)	58.89	60.41	59.64	48.62	45.52	47.02	51.7	46.04	48.71	59.25	58.09	59.67
Sequence-tagging												
Span-BART (Yan et al., 2021)	-	-	72.46	-	-	57.59	-	-	60.10	-	-	69.98
JET (Xu et al., 2020)	67.97	60.32	63.92	58.47	43.67	50.00	58.35	51.43	54.67	64.77	61.29	62.98
MRC based												
BMRC [†] (Chen et al., 2021a)	71.32	70.09	70.69	65.12	54.41	59.27	63.71	58.63	61.05	67.74	68.56	68.13
COM-MRC (Zhai et al., 2022)	<u>76.45</u>	69.67	72.89	64.73	56.09	60.09	68.50	59.74	63.65	<u>72.80</u>	70.85	71.79
Table-filling												
S ³ E ² (Chen et al., 2021b)	69.08	64.55	66.74	59.43	46.23	52.01	61.06	56.44	58.66	71.08	63.13	66.87
GTS (Wu et al., 2020a)	70.92	69.49	70.20	57.52	51.92	54.58	59.29	58.07	58.67	68.58	66.60	67.58
EMC-GCN (Chen et al., 2022)	71.85	72.12	71.78	61.46	<u>55.56</u>	58.32	59.89	61.05	60.38	65.08	71.66	68.18
BDTF (Zhang et al., 2022)	76.71	<u>74.01</u>	<u>75.33</u>	68.30	55.10	<u>60.99</u>	66.95	65.05	65.97	73.43	<u>73.64</u>	73.51
DGCNAP (Li et al., 2023)	71.83	68.77	70.26	66.46	54.34	58.74	62.03	57.18	59.49	69.39	72.20	70.77
LLM-based												
GPT 3.5 zero-shot	39.21	56.17	46.18	26.21	40.69	31.88	31.21	52.75	39.21	35.28	59.64	44.34
GPT 3.5 few-shots	44.73	58.87	50.84	29.63	37.69	33.18	37.27	56.42	44.89	43.15	60.75	50.46
GPT 4 zero-shot	27.34	37.13	31.49	16.50	24.41	19.69	25.60	39.22	30.98	28.39	43.64	35.79
GPT 4 few-shots	41.48	52.06	46.17	28.79	39.83	33.42	38.04	58.02	45.96	40.89	62.50	49.44
Ours												
ContrASTE	75.87	76.12	76.00	<u>67.45</u>	61.01	64.07	<u>66.84</u>	<u>64.08</u>	<u>65.43</u>	69.38	74.40	<u>71.80</u>

Table 7: Experimental results on \mathcal{D}_1 (Wu et al., 2020a). The best results are highlighted in bold, while the second best results are underlined.

Methods	Brief Introduction
Pipeline	
OTE-MTL (Zhang et al., 2020)	It proposes a multi-task learning framework including two parts: aspect and opinion tagging, along with word-level sentiment dependency parsing. This approach simultaneously extracts aspect and opinion terms while parsing sentiment dependencies using a biaffine scorer. Additionally, it employs triplet decoding based on the aforementioned outputs during inference to facilitate triplet extraction.
Li-unified-R+PD (Peng et al., 2020)	It proposes an unified tagging scheme, Li-unified-R, to assist target boundary detection. Two stacked LSTMs are employed to complete aspect-based sentiment prediction and the sequence labeling.
CMLA+C-GCN (Wang et al., 2017)	It facilitates triplet extraction by modelling the interaction between the aspects and opinions.
Two-satge (Peng et al., 2020)	It decomposes triplet extraction to two stages: 1) predicting unified aspect-sentiment and opinion tags; and 2) pairing the two results from stage one.
RI-NANTE+ (Dai and Song, 2019)	It adopts the same sentiment triplets extracting method as that of CMLA+, but it incorporates a novel LSTM-CRF mechanism and fusion rules to capture word dependencies within sentences.
Sequence-tagging	
Span-BART (Yan et al., 2021)	It redefines triplet extraction within an end-to-end framework by utilizing a sequence composed of pointer and sentiment class indexes. This is achieved by leveraging the pretrained sequence-to-sequence model BART to address ASTE.
JET (Xu et al., 2020)	It extracts triplets jointly by designing a position-aware sequence-tagging scheme to extract the triplets and capturing the rich interactions among the elements.
Seq2seq	
Dual-MRC (Mao et al., 2021)	It proposes a solution for ASTE by jointly training two BERT-MRC models with parameters sharing.
BMRC (Chen et al., 2021a)	It introduces a bidirectional MRC (BMRC) framework for ASTE, employing three query types: non-restrictive extraction queries, restrictive extraction queries, and sentiment classification queries. The framework synergistically leverages two directions, one for sequential recognition of aspect-opinion-sentiment and the other for sequential recognition of opinion-aspects-sentiment expressions.
Table-filling	
GTS (Wu et al., 2020a)	It proposes a novel 2D tagging scheme to address ASTE in an end-to-end fashion only with one unified grid tagging task. It also devises an effective inference strategy on GTS that utilizes mutual indication between different opinion factors to achieve more accurate extraction.
Double-encoder (Jing et al., 2021)	It proposes a dual-encoder model that capitalizes on encoder sharing while emphasizing differences to enhance effectiveness. One of the encoders, referred to as the pair encoder, specifically concentrates on candidate aspect-opinion pair classification, while the original encoder retains its focus on sequence labeling.
S ³ E ² (Chen et al., 2021b)	It represents the semantic and syntactic relationships between word pairs, employs GNNs for encoding, and applies a more efficient inference strategy.
EMC-GCN (Chen et al., 2022)	It employs a biaffine attention module to embed ten types of relations within sentences, transforming the sentence into a multi-channel graph while incorporating various linguistic features to enhance performance. Additionally, the method introduces an effective strategy for refining word-pair representations, aiding in the determination of whether word pairs are a match or not.
LLM-based	
zero-shot	Performing aspect-based sentiment analysis using an LLM. The specific method involves inputting a prompted sentence and directly outputting the corresponding [A, O, S] triplets. An example of the text given to the LLM, with the prompt added, is as follows: "Perform aspect-based sentiment analysis on the provided text and return triplets as [Aspect, Opinion, Sentiment]. You only need to provide the triplets, no additional explanations are required. The provided text: {sentence}"
few-shots	Building upon the zero-shot method, a small number of examples from the training set are added to the prompted sentence: "Perform aspect-based sentiment analysis on the provided text and return triplets as [Aspect, Opinion, Sentiment]. For example: input: {train sentence} output: {train triplets}, ... (some other examples). You only need to provide the triplets, no additional explanations are required. The provided text: {sentence}"

Table 8: Baselines methods with brief introduction.

Sentence	Ground Truth	Predictions			
		GPT-3.5 zero-shot	GPT-3.5 few-shots	GPT-4 zero-shot	GPT-4 few-shots
It is a cozy place to go with a couple of friends.	[place, cozy, positive]	[place, cozy, positive], [friends, couple of, positive]	[place, cozy, positive]	[place, cozy, positive] [going with friends, a couple of, positive]	[place, cozy, positive]
This place has the best sushi in the city.	[sushi, best, positive]	[sushi, best, positive], [place, best, positive], [city, best, positive]	[sushi, best, positive]	[sushi, the best, positive], [place, This, neutral], [city, in the, neutral]	[sushi, best, positive]
Disappointingly, their wonderful Saketini has been taken off the bar menu.	[Saketini, wonderful, positive], [bar menu, Disappointingly, negative]	[Saketini, wonderful, positive], [bar menu, taken off, negative]	[Saketini, taken off, negative]	[Saketini, has been taken off the bar menu, negative], [bar menu, wonderful, positive]	[Saketini, wonderful, positive], [bar menu, taken off, negative]
I went in one day asking for a table for a group and was greeted by a very rude hostess.	[hostess, rude, negative]	[hostess, rude, negative]	[hostess, rude, negative]	[table for a group, asking, neutral], [hostess, very rude, negative]	[table for a group, asking, neutral], [hostess, very rude, negative]
But make sure you have enough room on your credit card as the bill will leave a big dent in your wallet.	[bill, big, negative]	[credit card, enough room, positive], [bill, big dent, negative], [wallet, big dent, negative]	[credit card, enough room, negative], [bill, big dent, negative], [wallet, big dent, negative]	[room on your credit card, enough, positive], [bill, will leave a big dent in your wallet, negative]	[bill, big dent, negative]

Table 9: In summary, there are several challenges observed in the performance of GPT models concerning triplets. Firstly, there is a prominent issue of "hard" matching, where GPT models tend to introduce additional modifiers or adverbs in the opinion component, leading to a lack of exact correspondence. Secondly, during zero-shot inference, GPT models tend to generate multiple predicted triplets, resulting in decreased precision. This behavior particularly hampers the precision of the model's predictions. Thirdly, inconsistencies arise in handling triplets involving structures such as [A, O1 and O2, S] and [A, O1, S], [A, O2, S]. This inconsistency is challenged by the evaluation metrics. dependence on annotation practices and conventions. Upon closer examination, the issues observed do not appear to be as pronounced as indicated by the evaluation metrics. Rather, they often manifest as cases where the general idea is correctly captured, but the precise format or phrasing does not align perfectly. Notably, the performance of GPT-4 deteriorates due to its occasional tendency to not merely "extract" fragments from sentences but to generate its own summarizations. Consequently, evaluating against triplets that originate solely from annotated sentences poses a challenge in achieving alignment. Furthermore, GPT-4 exhibits a proclivity for extracting longer sequences of words as aspects or opinions, while GPT-3.5 tends to produce shorter sequences that better conform to typical annotation scenarios.